

Titanic_Jupyter

March 6, 2021

```
[11]: # -*- Project -*-
      """
      Created on Sun Dec 20 12:09:18 2020

      @author: Georg Z.
      """

      #Read Data
      import pandas as pd
      titanic = pd.read_csv("train.csv", sep = ',')    #train data
      test = pd.read_csv("test.csv", sep = ',')      #data to be predicted

[12]: #####
      #Model Meta-Parmeters
      cv=2
      polynom_degree = 2

[13]: #Describe Data
      titanic.head()
      titanic.keys()
      titanic.Survived
      titanic.shape
      titanic.dtypes

[13]: PassengerId      int64
      Survived         int64
      Pclass           int64
      Name             object
      Sex              object
      Age             float64
      SibSp            int64
      Parch            int64
      Ticket           object
      Fare            float64
      Cabin            object
      Embarked         object
      dtype: object
```

```
[14]: #show non-numeric object columns
obj_titanic = titanic.select_dtypes(include=['object']).copy()
obj_titanic.head()

sum_stat = round(titanic.describe(),1)
print(sum_stat.to_latex())
print(titanic)

print(titanic.dtypes.to_latex())
```

```
\begin{tabular}{lrrrrrrrr}
\toprule
{} & PassengerId & Survived & Pclass & Age & SibSp & Parch & Fare \\
\midrule
count & 891.0 & 891.0 & 891.0 & 714.0 & 891.0 & 891.0 & 891.0 \\
\\
mean & 446.0 & 0.4 & 2.3 & 29.7 & 0.5 & 0.4 & 32.2 \\
\\
std & 257.4 & 0.5 & 0.8 & 14.5 & 1.1 & 0.8 & 49.7 \\
\\
min & 1.0 & 0.0 & 1.0 & 0.4 & 0.0 & 0.0 & 0.0 \\
\\
25\% & 223.5 & 0.0 & 2.0 & 20.1 & 0.0 & 0.0 & 7.9 \\
\\
50\% & 446.0 & 0.0 & 3.0 & 28.0 & 0.0 & 0.0 & 14.5 \\
\\
75\% & 668.5 & 1.0 & 3.0 & 38.0 & 1.0 & 0.0 & 31.0 \\
\\
max & 891.0 & 1.0 & 3.0 & 80.0 & 8.0 & 6.0 & 512.3 \\
\\
\bottomrule
\end{tabular}
```

```

PassengerId Survived Pclass \
0           1         0      3
1           2         1      1
2           3         1      3
3           4         1      1
4           5         0      3
..          ...         ...   ...
886         887         0      2
887         888         1      1
888         889         0      3
889         890         1      1
890         891         0      3
```

```
Name Sex Age SibSp \
```

0		Braund, Mr. Owen Harris	male	22.0	1
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2		Heikkinen, Miss. Laina	female	26.0	0
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4		Allen, Mr. William Henry	male	35.0	0
..		
886		Montvila, Rev. Juozas	male	27.0	0
887		Graham, Miss. Margaret Edith	female	19.0	0
888	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	
889		Behr, Mr. Karl Howell	male	26.0	0
890		Dooley, Mr. Patrick	male	32.0	0

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
..	
886	0	211536	13.0000	NaN	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	NaN	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	NaN	Q

[891 rows x 12 columns]

```

\begin{tabular}{ll}
\toprule
{} & 0 \\
\midrule
PassengerId & int64 \\
Survived & int64 \\
Pclass & int64 \\
Name & object \\
Sex & object \\
Age & float64 \\
SibSp & int64 \\
Parch & int64 \\
Ticket & object \\
Fare & float64 \\
Cabin & object \\
Embarked & object \\
\bottomrule
\end{tabular}

```

```
[15]: #mean of survived is more towards 0, which means that sample is imbalanced
#more people did not survive, than those who did. Eventually have to use
#stratified CV
```

```
titanic["Survived"].sum()

#Women/Men
titanic.loc[titanic.Sex == "female"]
titanic.loc[titanic.Sex == "male"]

#Women/Men survival rate
titanic.loc[titanic.Sex == "female"][titanic.Survived == 1]
titanic.loc[titanic.Sex == "male"][titanic.Survived == 1]

#round(sum(men)/len(men)*100,1)
#round(sum(women)/len(women)*100,1)
#20 of men survived, 75 of women survived
```

<ipython-input-15-167ecdf373d8>:12: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```
titanic.loc[titanic.Sex == "female"][titanic.Survived == 1]
```

<ipython-input-15-167ecdf373d8>:13: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```
titanic.loc[titanic.Sex == "male"][titanic.Survived == 1]
```

```
[15]:
```

	PassengerId	Survived	Pclass	Name	Sex	\
17	18	1	2	Williams, Mr. Charles Eugene	male	
21	22	1	2	Beesley, Mr. Lawrence	male	
23	24	1	1	Sloper, Mr. William Thompson	male	
36	37	1	3	Mamee, Mr. Hanna	male	
55	56	1	1	Woolner, Mr. Hugh	male	
..	
838	839	1	3	Chip, Mr. Chang	male	
839	840	1	1	Marechal, Mr. Pierre	male	
857	858	1	1	Daly, Mr. Peter Denis	male	
869	870	1	3	Johnson, Master. Harold Theodor	male	
889	890	1	1	Behr, Mr. Karl Howell	male	

	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
17	NaN	0	0	244373	13.0000	NaN	S
21	34.0	0	0	248698	13.0000	D56	S
23	28.0	0	0	113788	35.5000	A6	S
36	NaN	0	0	2677	7.2292	NaN	C
55	NaN	0	0	19947	35.5000	C52	S
..
838	32.0	0	0	1601	56.4958	NaN	S
839	NaN	0	0	11774	29.7000	C47	C


857	51.0	0	0	113055	26.5500	E17	S
869	4.0	1	1	347742	11.1333	NaN	S
889	26.0	0	0	111369	30.0000	C148	C

[109 rows x 12 columns]

```
[16]: #ticket class and survival rate
passengers = [len(titanic.loc[titanic.Pclass == 1]), len(titanic.loc[titanic.
    ↪Pclass == 2]),
len(titanic.loc[titanic.Pclass == 3])]
```

```
[17]: import matplotlib.pyplot as plt

fig = plt.figure()
ax = fig.add_axes([0,0,1,1])
classes = ['1st Class', '2nd Class', '3rd Class']
ax.bar(classes,passengers, color=['black', 'grey', 'silver'])
plt.title('Passengers by Class')
plt.xlabel('Class')
plt.ylabel('Passenger Count')
plt.show()
```



```
[18]: passengers_surv_rate = [len(titanic.loc[titanic.Pclass == 1][titanic.Survived_
    ↪== 1])/len(titanic.loc[titanic.Pclass == 1])*100,
    len(titanic.loc[titanic.Pclass == 2][titanic.Survived == 1])/
    ↪len(titanic.loc[titanic.Pclass == 2])*100,
    len(titanic.loc[titanic.Pclass == 3][titanic.Survived == 1])/
    ↪len(titanic.loc[titanic.Pclass == 3])*100]

fig2 = plt.figure()
ax = fig2.add_axes([0,0,1,1])
classes = ['1st Class', '2nd Class', '3rd Class']
ax.bar(classes,passengers_surv_rate, color=['black', 'grey', 'silver'])
plt.title('% Survived by Class')
plt.xlabel('Class')
plt.ylabel('% survived')
plt.show()
```

<ipython-input-18-f83440d09da8>:1: UserWarning: Boolean Series key will be reindexed to match DataFrame index.


```
passengers_surv_rate = [len(titanic.loc[titanic.Pclass == 1][titanic.Survived
== 1])/len(titanic.loc[titanic.Pclass == 1])*100,
```

<ipython-input-18-f83440d09da8>:2: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```
len(titanic.loc[titanic.Pclass == 2][titanic.Survived ==
1])/len(titanic.loc[titanic.Pclass == 2])*100,
```

<ipython-input-18-f83440d09da8>:3: UserWarning: Boolean Series key will be reindexed to match DataFrame index.

```
len(titanic.loc[titanic.Pclass == 3][titanic.Survived ==
1])/len(titanic.loc[titanic.Pclass == 3])*100]
```



```
[19]: #age and survival
titanic["Age"].max()
titanic["Age"].min()
titanic["Age"].describe()

#print('Train Data Info')
#print(titanic.info())
```

```
[19]: count    714.000000
mean      29.699118
std       14.526497
min        0.420000
25%       20.125000
50%       28.000000
75%       38.000000
max       80.000000
Name: Age, dtype: float64
```

```
[20]: #####
#save passenger ids to assign them to predictions later
```

```
def save_id(data):
    id = data['PassengerId']
    return id

#Observe Missing Values
ids_test = save_id(test)
test.isna().sum()
test.isna().sum()
```

```
[20]: PassengerId      0
      Pclass         0
      Name           0
      Sex            0
      Age            86
      SibSp          0
      Parch          0
      Ticket         0
      Fare           1
      Cabin         327
      Embarked       0
      dtype: int64
```

```
[21]: #Pre-Process Data
      #observe missing values
      pd.isnull(titanic).sum()
      pd.isnull(test).sum()

      #create bins for fare
      # =====
      # def create_bins(x):
      #     range_value = [0,0,0,0]
      #     if 0 < x < 50:
      #         range_value[0] = 1
      #     elif 50 < x < 100:
      #         range_value[1] = 1
      #     elif 100 < x < 250:
      #         range_value[2] = 1
      #     else:
      #         range_value[3] = 1
      #     return range_value
      # =====

      def pre_process (data):
          #Drop Variables that don't contribute to outcome
          data = data.drop(columns=["PassengerId", "Name","Ticket","Cabin"])
          #only include rows in which embarked is not missing
          #titanic loses two observations, test loses none
```



```

    data = data.loc[data.Embarked.notna()]
    data = data.interpolate()

# =====
#     #bin age
#     import numpy as np
#     bins = [-1, 0, 5, 12, 18, 24, 35, 60, np.inf]
#     labels = ['Unknown', 'Baby', 'Child', 'Teenager', 'Student', 'Young
↳Adult', 'Adult', 'Senior']
#     data['AgeGroup'] = pd.cut(data["Age"], bins, labels = labels)
#     #bin fare
#     data["Fare_lt_50"], data["Fare_50_100"], data["Fare_100_250"],
↳data["Fare_gt_250"] = zip(*data["Fare"].map(create_bins))
#     data = data.drop(columns=["Fare"])
# =====

    return data

titanic = pre_process(titanic)
test = pre_process(test)
pd.isnull(titanic).sum()
pd.isnull(test).sum()

```

```

[21]: Pclass      0
      Sex         0
      Age         0
      SibSp       0
      Parch       0
      Fare        0
      Embarked    0
      dtype: int64

```

```

[22]: #####
      #One Hot Encoding
      def one_hot (data):
          from sklearn.preprocessing import OneHotEncoder
          enc = OneHotEncoder()

          #One hot encode Pclass
          Pclass = enc.fit_transform(data[["Pclass"]]).toarray()
          Pclass = pd.DataFrame(Pclass)
          Pclass.rename(columns={0: "1st Class", 1: "2nd Class", 2: "3rd Class"},
↳inplace=True)

          #One hot encode Sex
          sex = enc.fit_transform(data[["Sex"]]).toarray()
          sex = pd.DataFrame(sex)
          sex.rename(columns={0: "female", 1: "male"}, inplace=True)

```

```

#One hot encode Embarked
embarked = enc.fit_transform(data[["Embarked"]]).toarray()
embarked = pd.DataFrame(embarked)
embarked.rename(columns={0: "C", 1: "Q", 2: "S"}, inplace=True)

# =====
# #One hote encode Age
# age = enc.fit_transform(data[["AgeGroup"]]).toarray()
# age = pd.DataFrame(age)
# age.rename(columns={0: "Baby", 1: "Child", 2: "Teenager", 3: "Student", 4:
→ "Young Adult", 5: "Adult", 6: "Senior"}, inplace=True)
#
# =====

#reset index because pd.concat mismatches indexes and creates nans
data.reset_index(drop=True, inplace=True)
sex.reset_index(drop=True, inplace=True)

#Add recoded columns to data frame
data = pd.concat([data, pd.DataFrame(sex)], axis=1)
data = pd.concat([data, pd.DataFrame(embarked)], axis=1)
#data = pd.concat([data, pd.DataFrame(age)], axis=1)
data = pd.concat([data, pd.DataFrame(Pclass)], axis=1)

#Drop old columns
#data = data.drop(columns=["Sex", "Embarked", "AgeGroup", "Age", "Pclass"])
data = data.drop(columns=["Sex", "Embarked", "Pclass"])
return data

titanic = one_hot(titanic)
test = one_hot(test)

```

```

[23]: #last check for missing values: no missing values
pd.isnull(titanic).sum()
pd.isnull(test).sum()
titanic.shape
test.shape

```

[23]: (418, 12)

```

[24]: #####
#min max scaler chosen because all other are one hot encoded between 0,1
def scale_min_max (data):
    #Scale some features
    from sklearn.preprocessing import MinMaxScaler
    scaler = MinMaxScaler()
    scaler.fit(data[["SibSp"]])
    data[["SibSp"]] = scaler.transform(data[["SibSp"]])

```

```

#Parch (Parents)
scaler.fit(data[["Parch"]])
data[["Parch"]] = scaler.transform(data[["Parch"]])
#Scale Age
scaler.fit(data[["Age"]])
data[["Age"]] = scaler.transform(data[["Age"]])
#Scale Fare
scaler.fit(data[["Fare"]])
data[["Fare"]] = scaler.transform(data[["Fare"]])
return data

titanic = scale_min_max(titanic)
test = scale_min_max(test)

#look at summary statistics of scaled data
titanic.max(axis=0)
titanic.min(axis=0)

#explain .scale command(which scaler is it: Standard Scaler)
#not worried about scaling problems or outliers
#possibly outliers by Fare

sum_stat = round(titanic.describe(),1)
print(sum_stat)
print(sum_stat.to_latex())

print(titanic.dtypes.to_latex())

#import sys
#sys.exit("Stop here")
#####

```

	Survived	Age	SibSp	Parch	Fare	female	male	C	Q \
count	889.0	889.0	889.0	889.0	889.0	889.0	889.0	889.0	889.0
mean	0.4	0.4	0.1	0.1	0.1	0.4	0.6	0.2	0.1
std	0.5	0.2	0.1	0.1	0.1	0.5	0.5	0.4	0.3
min	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
25%	0.0	0.3	0.0	0.0	0.0	0.0	0.0	0.0	0.0
50%	0.0	0.4	0.0	0.0	0.0	0.0	1.0	0.0	0.0
75%	1.0	0.5	0.1	0.0	0.1	1.0	1.0	0.0	0.0
max	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0

	S	1st Class	2nd Class	3rd Class
count	889.0	889.0	889.0	889.0

```

mean      0.7      0.2      0.2      0.6
std       0.4      0.4      0.4      0.5
min       0.0      0.0      0.0      0.0
25%      0.0      0.0      0.0      0.0
50%      1.0      0.0      0.0      1.0
75%      1.0      0.0      0.0      1.0
max       1.0      1.0      1.0      1.0

```

```
\begin{tabular}{lrrrrrrrrrrrr}
```

```
\toprule
```

```

{} & Survived & Age & SibSp & Parch & Fare & female & male & C &
Q & S & 1st Class & 2nd Class & 3rd Class & \&

```

```
\midrule
```

```

count & 889.0 & 889.0 & 889.0 & 889.0 & 889.0 & 889.0 & 889.0 &
889.0 & 889.0 & 889.0 & 889.0 & 889.0 & 889.0 & \&
mean & 0.4 & 0.4 & 0.1 & 0.1 & 0.1 & 0.4 & 0.6 &
0.2 & 0.1 & 0.7 & 0.2 & 0.2 & 0.6 & \&
std & 0.5 & 0.2 & 0.1 & 0.1 & 0.1 & 0.5 & 0.5 &
0.4 & 0.3 & 0.4 & 0.4 & 0.4 & 0.5 & \&
min & 0.0 & 0.0 & 0.0 & 0.0 & 0.0 & 0.0 & 0.0 &
0.0 & 0.0 & 0.0 & 0.0 & 0.0 & 0.0 & \&
25\% & 0.0 & 0.3 & 0.0 & 0.0 & 0.0 & 0.0 & 0.0 &
0.0 & 0.0 & 0.0 & 0.0 & 0.0 & 0.0 & \&
50\% & 0.0 & 0.4 & 0.0 & 0.0 & 0.0 & 0.0 & 1.0 &
0.0 & 0.0 & 1.0 & 0.0 & 0.0 & 1.0 & \&
75\% & 1.0 & 0.5 & 0.1 & 0.0 & 0.1 & 1.0 & 1.0 &
0.0 & 0.0 & 1.0 & 0.0 & 0.0 & 1.0 & \&
max & 1.0 & 1.0 & 1.0 & 1.0 & 1.0 & 1.0 & 1.0 &
1.0 & 1.0 & 1.0 & 1.0 & 1.0 & 1.0 & \&

```

```
\bottomrule
```

```
\end{tabular}
```

```
\begin{tabular}{ll}
```

```
\toprule
```

```
{ } & 0 & \&
```

```
\midrule
```

```

Survived & int64 & \&
Age & float64 & \&
SibSp & float64 & \&
Parch & float64 & \&
Fare & float64 & \&
female & float64 & \&
male & float64 & \&
C & float64 & \&
Q & float64 & \&
S & float64 & \&
1st Class & float64 & \&
2nd Class & float64 & \&
3rd Class & float64 & \&

```

```
\bottomrule
\end{tabular}
```

[25]: *#Higher-Order Polynomials*

```
def feat_engin (data):
    from sklearn.preprocessing import PolynomialFeatures
    poly_transformer = PolynomialFeatures(degree=polynom_degree,
    ↪interaction_only=True, include_bias=True)
    data_poly = poly_transformer.fit_transform(data.loc[:, data.columns !=
    ↪'Survived'])

    #Add "Survived" column
    data_poly = pd.DataFrame(data_poly)
    if len(data) > len(test):      #only the train set receives a target column
        data_poly = pd.concat([data_poly, pd.DataFrame(data["Survived"])],
    ↪axis=1)
    else:
        data_poly = data_poly #the test set does not have a target column
    ↪"Survived"
    return data_poly

titanic_poly = feat_engin(titanic)
test_poly = feat_engin(test)

#####
```

[26]: *#Grid Search*

```
#####
#Import Model APIs
import warnings
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.svm import SVC
#from sklearn.calibration import CalibratedClassifierCV
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import VotingClassifier

#Define Models and Parameters
#kNN
param_knn = {'n_neighbors': [3,5,11,19],
             'weights': ['uniform', 'distance'],
             'metric': ['euclidean', 'manhattan']}
```

```

    }
kNN = KNeighborsClassifier()

#SVC
param_svc = {'C': [0.001, 0.01, 0.1, 1, 10, 100],
             'gamma': [0.001, 0.01, 0.1, 1, 10, 100],
             'random_state' : [1],
             'probability' : [True]
            }
svc = SVC()

#LogReg
import numpy as np
param_logreg = [
    {'penalty' : ['l1', 'l2', 'elasticnet', 'none'],
     'C' : np.logspace(-4, 4, 5),
     'max_iter' : [1000],
     'random_state' : [1]
    }
]
logreg = LogisticRegression()

#Forest
#left to default parameters because it takes too long
param_forest = [
    {'random_state' : [1],
     'min_samples_leaf': [1, 2],
     'min_samples_split': [2, 5]
    }
]
#empty because defaults work very good
forest = RandomForestClassifier()

#Grad Boosting
#left to default parameters because it takes too long
param_grbt = [
    {'random_state' : [1], 'learning_rate':[0.1,0.01,0.001],
     'max_depth':[2,3,4,5]
    }
]
#empty because defaults work very good
grbt = GradientBoostingClassifier()

#MLP
param_mlp = [
    {
        'activation': ['relu'],

```

```

        'solver': ['adam'],
        'alpha': [0.0001, 0.05],
        'learning_rate': ['constant', 'adaptive'],
        'max_iter': [1000]
    }
]
mlp = MLPClassifier()

#Voting Classifier
estimators=[("knn", knn), ("log_reg", logreg),
            ("rf", forest), ("gbdt", gbdt), ("svc", svc), ("mlp", mlp)]
#create our voting classifier, inputting all models
ensemble = VotingClassifier(estimators, voting="hard")
param_ensemble = [
    {
    }
]

```

```

[27]: #Define Tuning Function
def tune(model, param_grid):
    grid_search = GridSearchCV(model, param_grid, cv=cv, verbose=False)
    return grid_search

```

```

[28]: #Define Scoring Function
def score(data, model_to_tune, parameters):
    with warnings.catch_warnings():
        warnings.simplefilter("ignore")
        #assess train and test fit of test set provided by Kaggle
        from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(data.
↳drop(columns=["Survived"]),
                                                    data['Survived'],
↳stratify=data['Survived'], random_state=1)
        from sklearn.model_selection import cross_val_score
        from sklearn.model_selection import StratifiedKFold
        Stkfold = StratifiedKFold(n_splits=3, shuffle=True)
        model = tune(model_to_tune, parameters)
        train = round(cross_val_score(model, X_train, y_train, cv=Stkfold).
↳mean(),2)
        test = round(cross_val_score(model, X_test, y_test, cv=Stkfold).
↳mean(),2)
        train_txt = "{model_to_tune} Training accuracy:"
        test_txt = "{model_to_tune} Testing accuracy:"
        print(train_txt.format(model_to_tune = str(model_to_tune)[:10]), train)
        print(test_txt.format(model_to_tune = str(model_to_tune)[:10]), test)
    return

```

```
[29]: #Model Evaluation
#####

#Fit and Score Models using Accuracy
print(" ")
print(" ")
print("Simple Model: Parts")
print(" ")
score(titanic, svc, param_svc)
score(titanic, knn, param_knn)
score(titanic, logreg, param_logreg)
score(titanic, forest, param_forest)
score(titanic, grbt, param_grbt)
score(titanic, mlp, param_mlp)
print(" ")
print("Simple Model Combined:")
score(titanic, ensemble, param_ensemble)
print(" ")
```

Simple Model: Parts

SVC() Training accuracy: 0.79
 SVC() Testing accuracy: 0.82
 KNeighbors Training accuracy: 0.76
 KNeighbors Testing accuracy: 0.81
 LogisticRe Training accuracy: 0.78
 LogisticRe Testing accuracy: 0.78
 RandomForest Training accuracy: 0.8
 RandomForest Testing accuracy: 0.78
 GradientBo Training accuracy: 0.79
 GradientBo Testing accuracy: 0.78
 MLPClassif Training accuracy: 0.78
 MLPClassif Testing accuracy: 0.82

Simple Model Combined:
 VotingClas Training accuracy: 0.79
 VotingClas Testing accuracy: 0.82

```
[30]: #####
#Fit on train set predict on test set and evaluate predictions with CMs

from sklearn.model_selection import train_test_split
```



```

X_train, X_test, y_train, y_test = train_test_split(titanic.
↳drop(columns=["Survived"]), titanic['Survived'],
↳stratify=titanic['Survived'], random_state=1)

def eval(model, parameters):
    #predict test set, resulting from split
    #from sklearn.model_selection import train_test_split
    #X_train, X_test, y_train, y_test = train_test_split(data.
↳drop(columns=["Survived"]), data['Survived'], stratify=data['Survived'],
↳random_state=1)
    #tuned_model = tune(model, parameters)

    fitted_model = model.fit(X_train, y_train)
    predictions = fitted_model.predict(X_test)

    #Confusion Matrix
    from sklearn.metrics import confusion_matrix
    print('Confusion matrix:')
    mat1 = confusion_matrix(y_test, predictions)
    print(mat1)

    #Sensitivity
    #from sklearn.metrics import recall_score
    #print(recall_score(y_test, predictions))

    sens = mat1[1,1]/(mat1[1,0]+mat1[1,1])
    print('Sensitivity:', round(sens,2))
    spec = mat1[0,0]/(mat1[0,1]+mat1[0,0])
    print('Specificity:', round(spec,2))
    return predictions

```

```

[31]: print(" ")
print("Ensemble")
pred_ensemble = eval(ensemble, param_ensemble)
print(" ")
print("Random Forest")
pred_forest = eval(forest, param_forest)
print(" ")
print("Log Reg")
pred_logreg = eval(logreg, param_logreg)
print(" ")
print("KNN")
pred_kNN = eval(kNN, param_knn)
print(" ")
print("Grbt")
pred_grbt = eval(grbt, param_grbt)

```

```
print(" ")
print("MLP")
pred_mlp = eval(mlp, param_mlp)
print(" ")
print("SVC")
pred_svc = eval(svc, param_svc)
```

Ensemble

C:\Users\zhele\anaconda3\lib\site-

packages\sklearn\normal_network_multilayer_perceptron.py:582:

ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

Confusion matrix:

```
[[129  9]
 [ 27 58]]
```

Sensitivity: 0.68

Specificity: 0.93

Random Forest

Confusion matrix:

```
[[120 18]
 [ 21 64]]
```

Sensitivity: 0.75

Specificity: 0.87

Log Reg

Confusion matrix:

```
[[119 19]
 [ 24 61]]
```

Sensitivity: 0.72

Specificity: 0.86

KNN

Confusion matrix:

```
[[124 14]
 [ 28 57]]
```

Sensitivity: 0.67

Specificity: 0.9

Grbt

Confusion matrix:

```
[[119 19]
 [ 26 59]]
```

Sensitivity: 0.69

Specificity: 0.86

MLP

Confusion matrix:

```
[[126  12]
 [ 22  63]]
```

Sensitivity: 0.74

Specificity: 0.91

SVC

Confusion matrix:

```
[[132   6]
 [ 32  53]]
```

Sensitivity: 0.62

Specificity: 0.96

C:\Users\zhele\anaconda3\lib\site-

packages\sklearn\normal_network_multilayer_perceptron.py:582:

ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

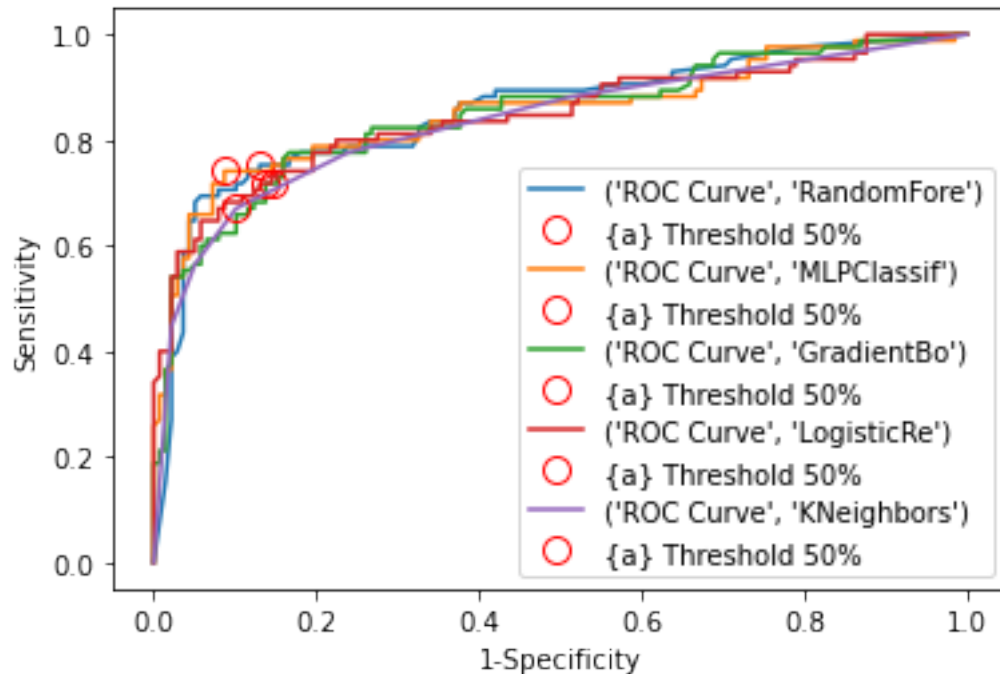
```
[32]: #Optimize Discrimination Threshold = optimize specificits/sensitivity trade-off

#ROC Curve
def roc(models):
    from sklearn.metrics import roc_curve
    for a in models:
        fpr, tpr, thresholds = roc_curve(y_test, a.predict_proba(X_test)[: , 1])
        label = ("ROC Curve", str(a)[:10])
        plt.plot(fpr, tpr, label=label)
        plt.xlabel("1-Specificity")
        plt.ylabel("Sensitivity")
        default = np.argmin(np.abs(thresholds - 0.5))
        plt.plot(fpr[default], tpr[default], 'o', markersize=10, label="{a}_
↪Threshold 50%", fillstyle="none", c='r')

        plt.legend(loc=4)
    return

models1=[forest, mlp]
roc(models1)
models2=[grbt, logreg, kNN]
roc(models2)
print("SVC has no Proba Eval Possibility")
```

SVC has no Proba Eval Possibility



[33]: *#Summary: Models are in good shape: no need for adjustment
 #the Thresholds of the two and the three curves are all on the 45 degree axis
 #therefore there is no opportunity for improvement, because Threshold on the
 → flat side of the curve
 #if try to improve will sacrifice specificity
 #lower prediction accuracy is due to other reasons than to Threshold selection
 #one can see that MLP and RF are the best models
 #perfect model ROC is as high sensitivity and as high specificity as possible*

[34]: *#Because ROC curves cross it is not possible to rank the models
 #therefore one can use a single evaluation measure to rankt the models
 #AUC is the integram of ROC
 #The higher it is, the better a model performs in terms of both sensitivity/
 → specificity
 #ranges between 0.5 and 1
 #random guessing would be an AUC of 0.5
 #AUC is insensitive to class proportions, due class inbalance
 #
 #AUC*

```
def auc(models):
    from sklearn.metrics import roc_auc_score
    for a in models:
        auc = roc_auc_score(y_test, a.predict_proba(X_test)[: , 1])
        print("AUC for", str(a)[:10], round(auc,2))
```

```

    return

models=[forest, grbt, logreg, kNN, mlp]
auc(models)
#random forest is the best model

```

```

AUC for RandomForest 0.85
AUC for GradientBo 0.85
AUC for LogisticRe 0.85
AUC for KNeighbors 0.84
AUC for MLPClassif 0.85

```

```

[35]: #Precision
def prec(predictions, models):
    from sklearn.metrics import precision_score
    for a, b in zip(predictions, models):
        print('Precision:', b)
        print(round(precision_score(y_test, a),2))
    return

predictions = [pred_ensemble, pred_logreg, pred_svc, pred_mlp, pred_forest,
    ↪pred_grbt]
prec(predictions, models)

```

```

Precision: RandomForestClassifier()
0.87
Precision: GradientBoostingClassifier()
0.76
Precision: LogisticRegression()
0.9
Precision: KNeighborsClassifier()
0.84
Precision: MLPClassifier()
0.78

```

```

[36]: #F Score
#sensitivity and precision combined
def f_score(predictions, models):
    from sklearn.metrics import f1_score
    for a, b in zip(predictions, models):
        print('Precision:', b)
        print(round(f1_score(y_test, a),2))
    return

f_score(predictions, models)

```

```

Precision: RandomForestClassifier()
0.76

```

```
Precision: GradientBoostingClassifier()
0.74
Precision: LogisticRegression()
0.74
Precision: KNeighborsClassifier()
0.79
Precision: MLPClassifier()
0.77
```

```
[37]: #Fit Complex Model
# a complex model of squares power does not improve performance
# =====
# print("Complex Model: Parts")
# print(" ")
# score(titanic_poly, svc, param_svc)
# score(titanic_poly, knn, param_knn)
# score(titanic_poly, logreg, param_logreg)
# score(titanic_poly, forest, param_forest)
# score(titanic_poly, grbt, param_grbt)
# score(titanic_poly, mlp, param_mlp)
# print(" ")
# print("Complex Model Combined:")
# score(titanic_poly, ensemble, param_ensemble)
# =====

#import sys
```

```
[38]: #####
print(" ")
print("Predict for Kaggle Competition:")
#predict "Survived" for test set provided by Kaggle
print(" ")
print("submission file is saved in the workign directory in csv format")
print("a submission based on the simple and one based on the complex model")

def predict_kaggle_test_set(data_to_predict, data_to_fit_on, model, parameters):
    #full train set provided by kaggle is used
    #test set provided by kaggle is predicted
    #prediction is done using the voting classifier
    X_data = data_to_fit_on.drop(columns=["Survived"])
    y_data = data_to_fit_on['Survived']
    model_tuned = tune(model, parameters)
    fitted_model = model_tuned.fit(X_data, y_data)
    predictions = fitted_model.predict(data_to_predict)
    return predictions

#prediction based on simple model
```

```
predictions_simple = predict_kaggle_test_set(test, titanic, ensemble,
↳param_ensemble)
```

Predict for Kaggle Competition:

submission file is saved in the workign directory in csv format
a submission based on the simple and one based on the complex model

```
C:\Users\zhele\anaconda3\lib\site-
packages\sklearn\normal_network\_multilayer_perceptron.py:582:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    warnings.warn(
C:\Users\zhele\anaconda3\lib\site-
packages\sklearn\normal_network\_multilayer_perceptron.py:582:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    warnings.warn(
C:\Users\zhele\anaconda3\lib\site-
packages\sklearn\normal_network\_multilayer_perceptron.py:582:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    warnings.warn(
```

```
[39]: #prediction based on complex model
#predictions_complex = predict_kaggle_test_set(test_poly, titanic_poly,
↳ensemble)

def save_csv(predictions, name):
    #set the output as a dataframe and convert to csv file
    output = pd.DataFrame({ 'PassengerId' : ids_test, 'Survived': predictions })
    output.to_csv(f"{name}.csv", index=False)
    return

save_csv(predictions_simple, "submission_simple")
#save_csv(predictions_complex, "submission_complex")
```

```
[40]: #####
#Comments on Titanic:
#Problem: train test accuracy 87%, validation set 78%. Model
#doesn't generalize good. Now generalizes better.
#see how the classes are represented in train test, and hold out data
#is titanic set class-imbalanced? (may need other metric as accuracy score)
#yes

#Submission after correction of expectations using cv.
```

```
[41]: #####  
      #Graphs/Tables  
  
      #discriptive  
      #ticket class and survival rate  
      #age and survival  
      #####  
  
      #Inference  
      #use stats models to show coefficients of logistic and linear reg
```