

Tweets

March 6, 2021

0.1 Sentiment Analysis with Disaster Tweets (Kaggle Competition Notebook)

0.1.1 Author: Georg Zhelev

This natural language processing projekt looks to determine when a natural disaster is taking place based on tweets from Twitter. SciKit Learn's bag of words method is applied coupled with a logistic regression, the parameters of which are chosen using grid search. Because the sample is about 60:40 unbalanced, a stratified cross validation is used to adjust prediction expectations. This model serves as a baseline for comparison with more advanced prediction methods.

0.1.2 Mount Drive

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

0.1.3 Import Tools

```
[ ]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import GridSearchCV

import numpy as np
from google.colab import files
```

0.1.4 Load Data

```
[ ]: df = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Econ/Disaster Tweets/
↳train.csv")
```

“df” is the test data

0.1.5 Explore Data

```
[ ]: df.head()
```

```
[ ]:   id keyword  ...                               text target
0    1     NaN  ...  Our Deeds are the Reason of this #earthquake M...      1
1    4     NaN  ...                Forest fire near La Ronge Sask. Canada      1
2    5     NaN  ...  All residents asked to 'shelter in place' are ...      1
3    6     NaN  ...  13,000 people receive #wildfires evacuation or...      1
4    7     NaN  ...  Just got sent this photo from Ruby #Alaska as ...      1
```

```
[5 rows x 5 columns]
```

```
[ ]: df.shape
```

```
[ ]: (7613, 5)
```

1: positive(forest fire) 0: negative (no forest fire)

```
[ ]: round(df.target.mean(),2)
```

```
[ ]: 0.43
```

```
[ ]: df.target.value_counts(normalize=True)
```

```
[ ]: 0    0.57034
     1    0.42966
     Name: target, dtype: float64
```

More tweets are not about forest fires (57%)

```
[ ]: df.isna().any()
```

```
[ ]: id            False
     keyword       True
     location      True
     text          False
     target        False
     dtype: bool
```

```
[ ]: print("Missing:")
     print("keyword missing:", df.keyword.isna().sum())
     print("location missing:", df.location.isna().sum())
```

```
Missing:
keyword missing: 61
location missing: 2533
```

What to do with missing location?

- For now don't use it as an explanation variable (worse)
- Use it as explanation variable, but drop those who are missing (better)
- Impute it somehow (best)

0.1.6 Pre-Process

```
[ ]: df1=df.dropna()
```

```
[ ]: 7613-(2533+61)
```

```
[ ]: 5019
```

Not exactly 5080 rows, because some are both NA for Location and Keyword

```
[ ]: df1.isna().any()
```

```
[ ]: id          False
      keyword     False
      location    False
      text        False
      target      False
      dtype: bool
```

0.1.7 Split Data

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(df["text"], df["target"],
      ↳test_size=0.33, stratify=df['target'], random_state=1)
```

```
[ ]: print(y_train.value_counts(y_train[1]))
```

```
0    0.570392
1    0.429608
Name: target, dtype: float64
```

```
[ ]: print(y_test.value_counts(y_test[0]))
```

```
0    0.570235
1    0.429765
Name: target, dtype: float64
```

```
[ ]: print(y_train.shape)
      print(y_test.shape)
```

```
(5100,)
(2513,)
```

```
[ ]: print(X_train.shape)
```

```
(5100,)
```

```
[ ]: X_train.head()
```

```
[ ]: 7360      'My Fifty Online Dates and why I'm still singl...
      826      Tomorrow's Announcement VODs http://t.co/cUbze...
      1877     @jaureguiswisdom lmao well i only know one and...
      7596     RT @LivingSafely: #NWS issues Severe #Thunders...
      4330     Criminals Who Hijack Lorries And Buses Arreste...
      Name: text, dtype: object
```

0.1.8 Bag of Words

```
[ ]: vect = CountVectorizer(min_df=3)
      vect.fit(X_train)
      X_train_bag = vect.transform(X_train)
      print('X_train:', repr(X_train_bag))
```

```
X_train: <5100x3330 sparse matrix of type '<class 'numpy.int64'>'
        with 59830 stored elements in Compressed Sparse Row format>
```

```
[ ]: features = vect.get_feature_names()
      print('First 10 features:', features[:10])
      print('Features 400 to 401', features[400:401])
      print('Every 100th feature:', features[::100])
```

```
First 10 features: ['00', '01', '02', '03', '04', '05', '06', '07', '08', '09']
Features 400 to 401 ['blanket']
Every 100th feature: ['00', 'aba', 'among', 'autumn', 'blanket', 'business',
'chile', 'cool', 'debate', 'dozens', 'eq', 'feat', 'friend', 'grown', 'hobbit',
'injured', 'kiernan', 'list', 'max', 'moral', 'north', 'page', 'plot', 'quake',
'rescue', 'salvador', 'shelter', 'southbound', 'sunday', 'thing', 'trfc',
'vancouver', 'whales', 'you']
```

0.1.9 Fit Model

```
[ ]: stkfold = StratifiedKFold(n_splits=5, shuffle=True)

      param_grid = {'C': [0.001, 0.01, 0.1, 1, 10]}

      grid = GridSearchCV(LogisticRegression(max_iter=100000), param_grid, cv=stkfold)
      grid.fit(X_train_bag, y_train)
```

```
[ ]: GridSearchCV(cv=StratifiedKFold(n_splits=5, random_state=None, shuffle=True),
                  error_score=nan,
```

```

estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                               fit_intercept=True,
                               intercept_scaling=1, l1_ratio=None,
                               max_iter=100000, multi_class='auto',
                               n_jobs=None, penalty='l2',
                               random_state=None, solver='lbfgs',
                               tol=0.0001, verbose=0,
                               warm_start=False),
iid='deprecated', n_jobs=None,
param_grid={'C': [0.001, 0.01, 0.1, 1, 10]},
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=None, verbose=0)

```

```

[ ]: print('Best cross-validation score:', grid.best_score_)
print('Best parameters:', grid.best_params_)
X_test_bag = vect.transform(X_test)
print(X_test_bag.shape)

print('Test accuracy', grid.score(X_test_bag, y_test))

```

```

Best cross-validation score: 0.7888235294117647
Best parameters: {'C': 0.1}
(2513, 3330)
Test accuracy 0.8038201352964585

```

Load Test Data

```

[ ]: test = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Econ/Disaster Tweets/
↳test.csv")

```

```

[ ]: test.shape

```

```

[ ]: (3263, 4)

```

```

[ ]: test.isna().any()

```

```

[ ]: id          False
keyword         True
location        True
text            False
dtype: bool

```

```

[ ]: print("keyword missing:", test.keyword.isna().sum())
print("location missing:", test.location.isna().sum())

```

```

keyword missing: 26
location missing: 1105

```

```
[ ]: test.head()
```

```
[ ]:      id keyword location      text
0    0      NaN      NaN  Just happened a terrible car crash
1    2      NaN      NaN  Heard about #earthquake is different cities, s...
2    3      NaN      NaN  there is a forest fire at spot pond, geese are...
3    9      NaN      NaN  Apocalypse lighting. #Spokane #wildfires
4   11      NaN      NaN  Typhoon Soudelor kills 28 in China and Taiwan
```

```
[ ]: id=test.id
```

```
[ ]: id.head()
```

```
[ ]: 0    0
     1    2
     2    3
     3    9
     4   11
     Name: id, dtype: int64
```

```
[ ]: X_test_k=test["text"]
     print(X_test_k.head())
```

```
0          Just happened a terrible car crash
1  Heard about #earthquake is different cities, s...
2  there is a forest fire at spot pond, geese are...
3          Apocalypse lighting. #Spokane #wildfires
4  Typhoon Soudelor kills 28 in China and Taiwan
     Name: text, dtype: object
```

here it is important to use the already fitted vocabulary using the count vectorizer and not fit a new one

```
[ ]: #vect.fit(X_test_k)
     X_test_k_bag = vect.transform(X_test_k)
     print('X_test_k:', repr(X_test_k_bag))
```

```
X_test_k: <3263x3330 sparse matrix of type '<class 'numpy.int64'>'
         with 36883 stored elements in Compressed Sparse Row format>
```

0.1.10 Make Predictions

```
[ ]: X_test_k_bag.shape
```

```
[ ]: (3263, 3330)
```

```
[ ]: X_train_bag.shape
```

```
[ ]: (5100, 3330)
```

```
[ ]: # predict results  
predictions = grid.predict(X_test_k_bag)
```

```
[ ]: predictions.shape
```

```
[ ]: (3263,)
```

```
[ ]: predictions.head()
```

```
[ ]: 0    0  
     1    0  
     2    1  
     3    0  
     4    1  
     Name: target, dtype: int64
```

```
[ ]: predictions = pd.Series(predictions, name="target")  
submission = pd.concat([pd.Series(id), predictions], axis = 1)
```

```
[ ]: submission.shape
```

```
[ ]: (3263, 2)
```

```
[ ]: submission.head()
```

```
[ ]:   id  target  
     0    0      0  
     1    2      0  
     2    3      1  
     3    9      0  
     4   11      1
```

```
[ ]: submission.to_csv("tweets_pred.csv", index=False)  
files.download("tweets_pred.csv")
```

<IPython.core.display.Javascript object>

<IPython.core.display.Javascript object>

```
[ ]: #Kaggle Competition Result  
  
#Accuracy on unlabeled test set: 0.78332
```